

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: data=pd.read_csv(r"C:\Users\LENOVO\Downloads\Advertising.csv")
data
```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [3]: data.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   TV          200 non-null    float64
 1   Radio        200 non-null    float64
 2   Newspaper    200 non-null    float64
 3   Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [5]: `data.tail()`

Out[5]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [6]: `data.describe`

Out[6]:


	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
..
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

[200 rows x 4 columns]>

In [7]: `import seaborn as sns
import matplotlib.pyplot as plt`

In [8]: `sns.pairplot(data,x_vars=['TV', 'Radio', 'Newspaper'],y_vars='Sales',height=7,aspect=0)`

Out[8]:



```
In [9]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression
```

```
In [10]: x=data.iloc[:,0:3]  
y=data.iloc[:, -1]
```

```
In [11]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20)  
a=LinearRegression()  
a.fit(x_train,y_train)
```

```
Out[11]:
```

```
  ▾ LinearRegression  
    LinearRegression()
```

```
In [12]: print(a.score(x_test,y_test))
```

```
0.9055191377503238
```

```
In [13]: from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import Lasso,Ridge
```

```
In [14]: print(a.coef_)
```

```
[0.05484129 0.09973974 0.00215296]
```


```
In [15]: features = data.columns[0:2]  
target = data.columns[-1]  
# x and y values  
x = data[features].values  
y = data[target].values  
#plot  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=17)  
print("The Dimension of x_train is {}".format(x_train.shape))  
print("The Dimension of x_test is {}".format(x_test.shape))  
#scale features  
scaler = StandardScaler()  
x_train = scaler.fit_transform(x_train)  
x_test = scaler.transform(x_test)
```

```
The Dimension of x_train is (140, 2)
```


```
The Dimension of x_test is (60, 2)
```

```
In [16]: plt.figure(figsize = (10, 10))
sns.heatmap(data.corr(), annot = True)
```

Out[16]: <Axes: >



```
In [17]: data.drop(columns = ["Radio", "Newspaper"], inplace = True)
#pairplot
sns.pairplot(data)
data.Sales = np.log(data.Sales)
```



```
In [18]: features = data.columns[0:2]
target = data.columns[-1]
#X and y values
X = data[features].values
y = data[target].values
#splot
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X_train is (140, 2)
 The dimension of X_test is (60, 2)

```
In [19]: #Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:


The train score for lr model is 1.0
The test score for lr model is 1.0

```
In [20]: #Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.990287139194161
The test score for ridge model is 0.9844266285141221

```
In [35]: plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color=
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```




```
In [32]: #Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls = lasso.score(X_train,y_train)
test_score_ls = lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0
 The test score for ls model is -0.0042092253233847465

```
In [22]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```


Out[22]: <Axes: >



```
In [23]: #Using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).fit(X_
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.999999343798134
 0.9999999152638072

```
In [31]: plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue'
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green'
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



```
In [36]: from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test))
```

The train score for ridge model is 0.999999999997627
The train score for ridge model is 0.9999999999962467

```
In [ ]:
```